

CS6240
Homework 2
Name: Rushabh Shah

GITHUB REPOSITORY: <https://github.ccs.neu.edu/rushabh0812/HW2>

Pseudo Codes:

No-Combiner:

Mapper extracts station id, type of the record and temperature from value v

Method map (Key k, Value v):

```
    If (recordType == 'TMAX' or recordType == 'TMIN'):  
        emit(stationID, (recordType, temperature, 1));
```

Method reduce(Key k, value [(recordType1, temperature1,1),.....]):

// key is the stationID

// value includes recordType, temperature and count

```
    maxCount = 0
```

```
    minCount = 0
```

```
    maxSum = 0
```

```
    minSum = 0
```

```
    minAvg = 0
```

```
    maxAvg = 0
```

```
    for each record in value v:
```

```
        if( v.recordType == 'TMAX'):
```

```
            maxSum += v.temperature
```

```
            maxCount += 1
```

```
        if( v.recordType == 'TMIN'):
```

```
            minSum += v.temperature
```

```
            minCount += 1
```

```
    minAvg = minSum/minCount
```

```
    maxAvg = maxSum/maxCount
```

```
    emit(stationID, minAvg, maxAvg)
```

Combiner:

//Mapper extracts station id, type of the record and temperature from value v

Method map (Key k, Value v):

```
    If (recordType == 'TMAX' or recordType == 'TMIN'):  
        emit(stationID, (recordType, temperature, 1));
```

Method combine(Key k, value [(recordType1, temperature1,1),.....]):

// key is the stationID

// value includes recordType, temperature and count

```
    maxCount = 0
```

```
    minCount = 0
```

```
    maxSum = 0
```

```
    minSum = 0
```

```
    for each record in value v:
```

```
        if( v.recordType == 'TMAX'):
```

```
            maxSum += v.temperature
```

```
            maxCount += 1
```

```
        if( v.recordType == 'TMIN'):
```

```
            minSum += v.temperature
```

```
            minCount += 1
```

```
    emit(stationID, ("TMAX", maxSum))
```

```
    emit(stationID, ("TMIN", minSum))
```

Method reduce(Key k, value [(recordType1, temperature1,1),.....]):

// key is the stationID

// value includes recordType, temperature and count

```
    maxCount = 0
```

```
    minCount = 0
```

```
    maxSum = 0
```

```
    minSum = 0
```

```
    minAvg = 0
```

```
    maxAvg = 0
```

```
    for each record in value v:
```

```
        if( v.recordType == 'TMAX'):
```

```
            maxSum += v.temperature
```

```

        maxCount += v.count
    if( v.recordType == 'TMIN'):
        minSum += v.temperature
        minCount += v.count

minAvg = minSum/minCount
maxAvg = maxSum/maxCount

emit(stationID, minAvg, maxAvg)

```

In Mapper Combiner:

```

class map{
    initialize():
    // initialize hashmaps
    Hmax , Hmin

    method map(Key k, Value v){
        sum= 0
        total = 0
        // extract station id, type of the record and temperature from value v
        if(recordType = 'TMAX'):
            sum = Hmax{stationID}.temperature + temperature
            total = Hmax{stationID}.count + 1
            Hmax.add(stationID, (sum, total))

        if(recordType = 'TMIN'):
            sum = Hmin{stationID}.temperature + temperature
            total = Hmin{stationID}.count + 1
            Hmax.add(stationID, (sum, total))

    Method cleanup():

        for each key in Hmax:
            emit(key, (Hmax{key}.sum + Hmax{key}.count))

        for each key in Hmin:
            emit(key, (Hmin{key}.sum + Hmin{key}.count))
    }
}

```

```

class reduce{

Method reduce(Key k, value [(recordType1, temperature1,1),..... ]):
// key is the stationID
// value includes recordType, temperature and count

    maxCount =0
    minCount = 0
    maxSum = 0
    minSum = 0
    minAvg = 0
    maxAvg = 0

    for each record in value v:
        if( v.recordType == "TMAX"):
            maxSum += v.temperature
            maxCount += 1
        if( v.recordType == "TMIN"):
            minSum += v.temperature
            minCount += 1

    minAvg = minSum/minCount
    maxAvg = maxSum/maxCount

    emit(stationID, minAvg, maxAvg)

}

```

Secondary Sort:

// map function converts the value into a key and value pair . The key is an object of CustomKey class

method map(Key k, value v) :

from value v extract station-id, year, record-type and temperature

if(record-type = "TMAX"):

emit((station-id,year), (year, 0,0,temperature,1))

if(record-type= "TMIN") :

emit((station-id,year),(year,temperature,1,0,0))

//The key is an object of type CustomKey. Custom key consists of two attributes. stationId and year.

class CustomKey {

stationId

year

method compareTo(Custom key k1 , Custom Key k2)

compare the stationId's .

if(same):

compare year

}

//The partitioner takes the key which is of type CustomKey and returns an appropriate partition based on stationId . All records with a particular station-id go to same reducer

method partitioner(key):

return hash(key.stationId)

//The combiner takes two parameters , key of the type CustomKey and list of values having same key

method combiner(key, values[(year, maxSum0,maxCount0, minSum0,minCount0),...]):

maxSum=0

minSum=0

maxCount=0

minCount=0

for each v in values:

maxSum += v.maxSum

minSum += v.minSum

maxCount += v.maxTempCount

minCount += v.minCount

emit(key, (year, maxSum, maxCount, minSum, minCount))

```
// The grouping comparator groups data by station id and all records having same stationId are sent to same reducer.
```

```
method customGroupComparator (Key k1, Key k2):  
//Key consists of station-id and year  
    compareValue = compare(k1.station-id, k2.station-id)  
    return compareValue
```

```
//The reduce function takes two parameters. A key of type CustomKey and list of values having same key . Each value contains five components. year , TMIN sum , TMIN, count , TMAX sum, TMAX count. The values are received in the increasing order of year
```

```
method reduce(key k, values[(year, minSum, maxSum, minCount, maxCount)] ):  
    maxSum=0  
    minSum=0  
    maxCount=0  
    minCount=0  
    year = key.year  
  
    for each in values:  
        if(v.year is not equal to year):  
            emit(key, (year, maxSum, maxCount, minSum, minCount))  
  
            maxSum=0  
            minSum=0  
            maxCount=0  
            minCount=0  
  
    maxSum += v.maxSum  
    minSum += v.minSum  
    maxCount += v.maxTempCount  
    minCount += v.minCount
```

The mapper emits records in the increasing order of stationId. This is achieved by overriding the inbuilt compareTo function. We defined our own custom compareTo function in CustomKey class . In this compareTo function , we first compare by stationId . If stationId's are equal, then we compare the years. Thus mapper emits records in increasing order of keys The grouping comparator the groups all the records having similar stationId and sends them to the reducer. Since we have defined our own custom comparator, the records in the reducer will be already present in the increasing order of year(In custom comparator , we compare by year

when stationId's are equal). Thus we make use of map reduce's sorting ability to prevent explicit sorting of values in reducer. This in-turn eliminates the need for complex data structures needed to sort.

Running time of the programs:

Program	Running Time 1 in seconds	Running Time 2 in seconds
No Combiner	108	104
Combiner	78	78
In Mapper Combiner	76	72
Secondary Sort	56	

Q. Was the Combiner called at all in program Combiner? Was it called more than once per Map task?

Ans: As we can see from the log files, the **Combiner was called** in the program for the two executions of the program Combiner.

First Execution:

Map-Reduce Framework

```
Map input records=31688662
Map output records=9213198
Map output bytes=267182742
Map output materialized bytes=3499100
Input split bytes=1560
Combine input records=9213198
Combine output records=548610
Reduce input groups=14723
Reduce shuffle bytes=3499100
Reduce input records=548610
Reduce output records=14723
Spilled Records=1097220
Shuffled Maps =180
Failed Shuffles=0
Merged Map outputs=180
```

Second Execution:

```

Total megabyte-milliseconds taken by all reducers
Map-Reduce Framework
  Map input records=31688662
  Map output records=9213198
  Map output bytes=267182742
  Map output materialized bytes=3499100
  Input split bytes=1560
  Combine input records=9213198
  Combine output records=548610
  Reduce input groups=14723
  Reduce shuffle bytes=3499100
  Reduce input records=548610
  Reduce output records=14723
  Spilled Records=1097220
  Shuffled Maps =180
  Failed Shuffles=0
  Merged Map outputs=180
  GC time elapsed (ms)=16172
-----
```

Q. Was the local aggregation effective in In-Mapper Combiner compared to No-Combiner?

Ans: **Yes**, the local aggregation in In-Mapper Combiner is effective compared to No-Combiner.

No Combiner Log records:

```

Total megabyte-milliseconds taken by all reducers
Map-Reduce Framework
  Map input records=31688662
  Map output records=9213198
  Map output bytes=230329950
  Map output materialized bytes=51438314
  Input split bytes=1560
  Combine input records=0
  Combine output records=0
  Reduce input groups=14723
  Reduce shuffle bytes=51438314
  Reduce input records=9213198
  Reduce output records=14723
  Spilled Records=18426396
  Shuffled Maps =180
  Failed Shuffles=0
```


In-Mapper Combiner Log records:

```
Total megabyte-milliseconds taken by all map tasks
Total megabyte-milliseconds taken by all reduce tasks
Map-Reduce Framework
  Map input records=31688662
  Map output records=233643
  Map output bytes=8411148
  Map output materialized bytes=4207356
  Input split bytes=1547
  Combine input records=0
  Combine output records=0
  Reduce input groups=14723
  Reduce shuffle bytes=4207356
  Reduce input records=233643
  Reduce output records=14723
  Spilled Records=467286
  Shuffled Maps =153
  Failed Shuffles=0
  Merged Map outputs=153
  GC time elapsed (ms)=17477
```

By looking at the log records for In-Mapper Combiner and No-Combiner, we can see that there is a considerable reduction in the Map output records of both the programs.

The time taken to execute In-Mapper program is less than No Combiner as the less records are passed between the mapper and reducer.